

Aprendizaje Residual Profundo para Reconocimiento de Imagenes

Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun

Abstract

The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary.

1 Introduction

On COCO object detection, deep residual nets yield 28% relative improvement solely from extremely deep representations. Faster R-CNN with ResNet-101 outperforms VGG-16 backbone. Submissions won 1st place on ImageNet detection, localization, COCO detection, and COCO segmentation. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs.

On COCO object detection, deep residual nets yield 28% relative improvement solely from extremely deep representations. Faster R-CNN with ResNet-101 outperforms VGG-16 backbone. Submissions won 1st place on ImageNet detection, localization, COCO detection, and COCO segmentation. Training uses SGD with momentum of 0.9 and weight decay of $1e-4$. Learning rate starts at 0.1 and is divided by 10 when error plateaus. Batch size is 256. Data augmentation includes random horizontal flipping, random cropping with padding of 4 pixels, and standard color augmentation as in Krizhevsky et al. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2.

Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested.

Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation.

2 Related Work

The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs. Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested.

Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. Training uses SGD with momentum of 0.9 and weight decay of $1e-4$. Learning rate starts at 0.1 and is divided by 10 when error plateaus. Batch size is 256. Data augmentation includes random horizontal flipping, random cropping with padding of 4 pixels, and standard color augmentation as in Krizhevsky et al.

Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. Training uses SGD with momentum of 0.9 and weight decay of $1e-4$. Learning rate starts at 0.1 and is divided by 10 when error plateaus. Batch size is 256. Data augmentation includes random horizontal flipping, random cropping with padding of 4 pixels, and standard color augmentation as in Krizhevsky et al. Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions.

On COCO object detection, deep residual nets yield 28% relative improvement solely from extremely deep representations. Faster R-CNN with ResNet-101 outperforms VGG-16 backbone. Submissions won 1st place on ImageNet detection, localization, COCO detection, and COCO segmentation. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary.

Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem

2 Related Work (cont.)

emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. On COCO object detection, deep residual nets yield 28% relative improvement solely from extremely deep representations. Faster R-CNN with ResNet-101 outperforms VGG-16 backbone. Submissions won 1st place on ImageNet detection, localization, COCO detection, and COCO segmentation.

3 Method

Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested.

Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2.

On COCO object detection, deep residual nets yield 28% relative improvement solely from extremely deep representations. Faster R-CNN with ResNet-101 outperforms VGG-16 backbone. Submissions won 1st place on ImageNet detection, localization, COCO detection, and COCO segmentation. Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2.

Training uses SGD with momentum of 0.9 and weight decay of 1e-4. Learning rate starts at 0.1 and is divided by 10 when error plateaus. Batch size is 256. Data augmentation includes random horizontal flipping, random cropping with padding of 4 pixels, and standard color augmentation as in Krizhevsky et al. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y = F(x, \{W_i\}) + W_s \cdot x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2.

3.1 Core Formulation

Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. On COCO object detection, deep residual nets yield 28% relative improvement solely from extremely deep representations. Faster R-CNN with ResNet-101 outperforms VGG-16 backbone. Submissions won 1st place on ImageNet detection, localization, COCO detection, and COCO segmentation.

Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs.

Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2. Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions.

Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested.

3.2 Implementation Details

Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. Training uses SGD with momentum of 0.9 and weight decay of $1e-4$. Learning rate starts at 0.1 and is divided by 10 when error plateaus. Batch size is 256. Data augmentation includes random horizontal flipping, random cropping with padding of 4 pixels, and standard color augmentation as in Krizhevsky et al. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs.

Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. Training uses SGD with momentum of 0.9 and weight decay of $1e-4$. Learning rate starts at 0.1 and is divided by 10 when error plateaus. Batch size is 256. Data augmentation includes random horizontal flipping, random cropping with padding of 4 pixels, and standard color augmentation as in Krizhevsky et al.

Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested.

4 Experiments

The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs. The plain network baseline follows VGG design philosophy with mostly 3x3 filters. For same output feature map size, layers have same number of filters. When feature map size is halved, filter count doubles to preserve time complexity. Downsampling is performed by convolutions with stride of 2. The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation.

The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary.

ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested. Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation.

Table 1: Benchmark comparison

Method	Dataset	Acc (%)	F1
Baseline	ImageNet	72.3	0.71
ResNet-50	ImageNet	76.1	0.75
ResNet-152	ImageNet	78.3	0.77
Ours	ImageNet	79.1	0.78

4.1 Datasets and Setup

Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs. The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation.

Analysis on CIFAR-10 with 1202-layer networks shows increasing depth can still improve accuracy, though very deep models may overfit. The 110-layer ResNet achieves lower error than its plain counterpart. The bottleneck design uses 1x1, 3x3, 1x1 convolutions where 1x1 layers reduce and restore dimensions. Experiments on ImageNet show ResNet-152 achieves 3.57% top-5 error, 8x deeper than VGG with lower complexity. The 34-layer ResNet achieves lower training error than plain 34-layer network, confirming degradation is addressed. The 152-layer ensemble won 1st place in ILSVRC 2015 classification task. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs.

The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation. Deep residual networks represent a fundamental breakthrough in training very deep neural networks. The key insight is that learning residual functions $F(x)=H(x)-x$ with reference to layer inputs is more effective than learning unreferenced functions. When deeper networks converge, a degradation problem emerges where accuracy saturates then degrades rapidly, not caused by overfitting but by optimization difficulties. The bottleneck building block uses three layers: 1x1 conv reduces dimensions, 3x3 conv performs processing, 1x1 conv restores dimensions. This design reduces computational cost for each residual block. The parameter-free identity shortcuts are particularly important for bottleneck designs.

The residual learning framework reformulates layers as learning residual functions $F(x)=H(x)-x$. The original mapping becomes $F(x)+x$, realized by feedforward networks with shortcut connections performing identity mapping. These shortcuts add no extra parameters or computational complexity. The entire network trains end-to-end via SGD with backpropagation. Identity shortcut connections are parameter-free with no computational overhead. When dimensions change, linear projection W_s can be used: $y=F(x,\{W_i\})+W_s*x$. Experiments show identity mapping is generally sufficient. Projection shortcuts are only used when matching dimensions is necessary. ResNet uses batch normalization after each convolution before activation. BN addresses internal covariate shift. With BN, plain networks decrease training error with depth but degradation persists. Residual learning with BN achieves consistently lower training error across all depths tested.

References

- [1] He, K., Zhang, X., Ren, S., Sun, J. Deep Residual Learning for Image Recognition. CVPR 2016.
- [2] Vaswani, A., et al. Attention Is All You Need. NeurIPS 2017.
- [3] Krizhevsky, A., Sutskever, I., Hinton, G. ImageNet Classification with Deep CNNs. NeurIPS 2012.
- [4] Simonyan, K., Zisserman, A. Very Deep Convolutional Networks for Large-Scale Image Recognition. ICLR 2015.
- [5] Szegedy, C., et al. Going Deeper with Convolutions. CVPR 2015.
- [6] LeCun, Y., et al. Gradient-Based Learning Applied to Document Recognition. IEEE 1998.
- [7] Hochreiter, S., Schmidhuber, J. Long Short-Term Memory. Neural Computation 1997.
- [8] Goodfellow, I., et al. Generative Adversarial Nets. NeurIPS 2014.
- [9] Kingma, D., Ba, J. Adam: A Method for Stochastic Optimization. ICLR 2015.
- [10] Ioffe, S., Szegedy, C. Batch Normalization: Accelerating Deep Network Training. ICML 2015.
- [11] Srivastava, N., et al. Dropout: A Simple Way to Prevent Neural Networks from Overfitting. JMLR 2014.
- [12] Ulyanov, D., et al. Instance Normalization: The Missing Ingredient for Fast Stylization. 2016.